

## Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

Final Year: High Performance Computing Lab 2022-23 Sem I

Class: Final Year (Computer Science and Engineering)

Year: 2022-23

Course: High Performance Computing Lab

### Practical No. 10

PRN: 2019BTECS00035

Title: **CUDA Programming**

1. Implement Matrix-matrix Multiplication using global memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

Link:

<https://github.com/Dnyaneshwar-dev/HPC-Lab-Sem1/blob/main/Assignment%2010/matrix-mult-global.cu>

<<<1,1>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
97.5	17603862851	1	17603862851.0	17603862851	17603862851	cudaDeviceSynchronize
2.5	456236713	3	152078904.3	17528	456144775	cudaMallocManaged
0.0	967954	3	322651.3	250554	415898	cudaFree
0.0	216125	1	216125.0	216125	216125	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	17603948659	1	17603948659.0	17603948659	17603948659	matrixMultiply(float*, float*, float*, int, int)

<<<2,4>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
90.0	2358090715	1	2358090715.0	2358090715	2358090715	cudaDeviceSynchronize
10.0	262263497	3	87421165.7	17304	262163694	cudaMallocManaged
0.0	934002	3	311334.0	244295	400852	cudaFree
0.0	63590	1	63590.0	63590	63590	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	2358080451	1	2358080451.0	2358080451	2358080451	matrixMultiply(float*, float*, float*, int, int)

<<<8,32>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
59.9	260104982	3	86701660.7	11851	260048905	cudaMallocManaged
39.8	172888012	1	172888012.0	172888012	172888012	cudaDeviceSynchronize
0.3	1367358	3	455786.0	362288	527116	cudaFree
0.0	52676	1	52676.0	52676	52676	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	172874260	1	172874260.0	172874260	172874260	matrixMultiply(float*, float*, float*, int, int)

2. Implement Matrix-Matrix Multiplication using shared memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

Link:

<https://github.com/Dnyaneshwar-dev/HPC-Lab-Sem1/blob/main/Assignment%2010/matrix-mult-shared.cu>

<<<1,1>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
98.6	17832282302	1	17832282302.0	17832282302	17832282302	cudaDeviceSynchronize
1.4	257490518	3	85830172.7	18556	257395847	cudaMallocManaged
0.0	1006901	3	335633.7	252898	464668	cudaFree
0.0	66076	1	66076.0	66076	66076	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	17832271711	1	17832271711.0	17832271711	17832271711	matrixMultiply(float*, float*, float*, int, int)

<<<2,4>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
90.2	2336132183	1	2336132183.0	2336132183	2336132183	cudaDeviceSynchronize
9.7	252101476	3	84033825.3	11986	252049153	cudaMallocManaged
0.0	948779	3	316259.7	235648	408710	cudaFree
0.0	45984	1	45984.0	45984	45984	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	2336118396	1	2336118396.0	2336118396	2336118396	matrixMultiply(float*, float*, float*, int, int)

<<<8,32>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
60.8	253054269	3	84351423.0	27415	252907993	cudaMallocManaged
38.5	160296268	1	160296268.0	160296268	160296268	cudaDeviceSynchronize
0.7	2857750	3	952583.3	513064	1499833	cudaFree
0.0	131710	1	131710.0	131710	131710	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	160289639	1	160289639.0	160289639	160289639	matrixMultiply(float*, float*, float*, int, int)

3. Implement Prefix sum using CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

Link:

<https://github.com/Dnyaneshwar-dev/HPC-Lab-Sem1/blob/main/Assignment%2010/prefix-sum.cu>

<<<1,1>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
56.0	272468280	4	68117070.0	6205	272352977	cudaMallocManaged
43.8	213044077	1	213044077.0	213044077	213044077	cudaDeviceSynchronize
0.2	867008	4	216752.0	18471	476186	cudaFree
0.1	342581	1	342581.0	342581	342581	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	213030832	1	213030832.0	213030832	213030832	prefixSum(float*, float*, float*, float*, int)

<<<1,32>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
94.1	419810388	4	104952597.0	14652	419567914	cudaMallocManaged
5.3	23787715	1	23787715.0	23787715	23787715	cudaDeviceSynchronize
0.4	1738409	4	434602.3	35540	910844	cudaFree
0.2	708542	1	708542.0	708542	708542	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	23780876	1	23780876.0	23780876	23780876	prefixSum(float*, float*, float*, float*, in t)

<<<1,64>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
92.6	259886563	4	64971640.8	5607	259762681	cudaMallocManaged
6.9	19435545	1	19435545.0	19435545	19435545	cudaDeviceSynchronize
0.3	963374	4	240843.5	26660	508137	cudaFree
0.1	365047	1	365047.0	365047	365047	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	19427656	1	19427656.0	19427656	19427656	prefixSum(float*, float*, float*, float*, in t)

CUDA Memory Operation Statistics (by time):

4. Implement 2D Convolution using shared memory using CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

Link:

<https://github.com/Dnyaneshwar-dev/HPC-Lab-Sem1/blob/main/Assignment%2010/convolution.cu>

Threads: 1

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.6	257739899	3	85913299.7	18701	257660598	cudaMallocManaged
0.3	815191	3	271730.3	42490	456648	cudaFree
0.0	120758	1	120758.0	120758	120758	cudaMemcpyToSymbol
0.0	37957	1	37957.0	37957	37957	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	16304415	1	16304415.0	16304415	16304415	convolution_2d(int*, int*, int)

CUDA Memory Operation Statistics (by time):

Threads: 32

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.6	257111316	3	85703772.0	18275	257031754	cudaMallocManaged
0.3	825387	3	275129.0	41599	483129	cudaFree
0.1	133246	1	133246.0	133246	133246	cudaMemcpyToSymbol
0.0	29602	1	29602.0	29602	29602	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	4014970	1	4014970.0	4014970	4014970	convolution_2d(int*, int*, int)

CUDA Memory Operation Statistics (by time):

Threads: 256

Exported successfully to  
/tmp/nsys-report-5002-6fe4-01b0-fe77.sqlite

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.8	263444022	3	87814674.0	17929	263374459	cudaMallocManaged
0.2	522316	3	174105.3	38338	315961	cudaFree
0.0	123902	1	123902.0	123902	123902	cudaMemcpyToSymbol
0.0	1289	1	1289.0	1289	1289	cudaLaunchKernel

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
100.0	2432	1	2432.0	2432	2432	[CUDA memcpy DtoD]

